LINGUAFORCE: Benchmarking Agentive Force of Discourse in Dialogues via Unified Psychological Dimensions

First Author, Second Author, Third Author
School of Computer Science, Example University
Example City, Country
{first, second, third}@example.edu

Abstract—Discourse that drives, guides, or restricts a listener's behavior—orders, threats, inducements, authority pressure, guilt-tripping, peer pressure, gaslighting, and verbal abuse—shares a common perlocutionary mechanism, yet existing NLP benchmarks study each strategy in isolation with task-specific labels and incompatible measurement scales. We propose LINGUAFORCE, a unified benchmark for the agentive force of discourse in real-world multi-turn dialogues. We design a three-layer framework: (i) a strategy taxonomy of four families and fifteen types that integrates compliance-gaining theories and speech-act theory; (ii) a universal inventory of seven psychological dimensions—directive force, option constraint, normative pressure, emotional pressure, deceptiveness, toxicity, and explicitness—grounded in pragmatics and persuasion research, which specializes to the obligation, constraint, value-judgement, and toxicity channels previously studied for moral coercion; and (iii) a linguistic-realization layer linking dimensions to surface cues. We construct LINGUAFORCE, a new dataset of 3,432 real-world multi-turn dialogues with dialogue-level annotations of intensity and the seven dimensions, re-annotated under the unified scheme. Annotation reliability is verified through consistency with the existing labels on the shared dialogues, and the benchmark supports zero-shot recognition of unseen manipulation types. We define three benchmark tasks, a two-stage dimension-aware pipeline, baseline models, and a cross-domain evaluation covering MentalManip, MultiManip, TalkDown, and ToxicChat.

Index Terms—Agentive force, perlocutionary pressure, computational social science, psycholinguistic framework, large language models, manipulation detection.

I. Introduction

Language does more than describe the world: utterances routinely act on their listeners, driving, guiding, or restricting behavior. This perlocutionary pressure appears along a broad spectrum of everyday strategies—a polite request that trades on reciprocity, an offer that reframes a decision, an authority figure's direct command, a threat or warning, social proof and conformity pressure, an emotional appeal, and covert tactics such as gaslighting. These strategies differ in surface form and intent, yet they share a common function: they operate on the listener's decision space, expanding or constraining the options that appear available. We call this capacity the agentive force of discourse.

Current NLP research fragments this spectrum across disconnected communities. Toxicity detection targets explicit aggression but misses polite or altruistic manipulation [1], [2]. Persuasion technique detection focuses on news articles and rhetorical framing [3]. Psychological-manipulation datasets cover a narrow subset of covert strategies such as gaslighting and intimidation and rarely include benign or transparent requests as controls [4]–[6]. Compliance-gaining research in communication studies provides rich taxonomies of persuasive strategies [7]–[9] but has not been operationalized as a unified NLP benchmark with continuous, comparable measurements. As a result, no existing resource answers three questions that matter for both science and safety: (i) what common structure underlies strategies as different as a polite request, a threat, and gaslighting; (ii) can a single measurement scale make them comparable; and (iii) can a model trained on a unified spectrum recognize manipulation forms it has never seen?

Only one point on this spectrum has been studied at scale—moral coercion, in which a speaker exploits shared ethical norms and levies preemptive moral labels such as "selfish" or "apathetic"—where decomposing the phenomenon into four psychological dimensions (obligation, constraint, value judgement, toxicity) was shown to yield more robust detection than direct classification [10], and it remains an open question whether such dimension-based decomposition transfers to the wider spectrum of linguistic influence.

To address these gaps, we introduce LINGUAFORCE, a benchmark design for the agentive force of discourse in real-world multi-turn dialogues. Our contributions are threefold:

1) A unified three-layer framework. We integrate compliance-gaining theory, speech-act theory, and persuasion-technique taxonomies into a strategy taxonomy of four families and fifteen types (Section III), a universal inventory of seven psychological dimensions (Section III-B), and a linguistic-realization layer that links each dimension to observable surface cues. The framework specializes to the obligation, constraint, value-judgement, and toxicity channels previously studied for one special case of the spectrum (moral coercion) as a consistency check.

2) A new dataset and benchmark. We construct LINGUAFORCE, a 3,432-dialogue dataset with dialogue-level annotations of intensity and seven psychological dimensions; annotation reliability is verified through consistency with existing labels on shared dialogues (Section IV).
3) An experimental protocol. We define three tasks (detection, type recognition, and intensity prediction), a two-stage dimension-aware pipeline, a zero-shot evaluation against a direct baseline, and research questions on dimension structure and zero-shot generalization to unseen manipulation types (Section V).

II. Related Work

A. Moral Coercion and Dimension Decoupling

The most direct prior evidence for dimensional analysis of coercion comes from a 4,700-dialogue study that annotates coercion intensity on a 0–5 scale and defines four psychological dimensions—obligation (OB), constraint (CS), value judgement (VJ), and toxicity (TX)—which a two-stage pipeline uses as structural anchors [10]. Its statistical analyses show that moral coercion overlaps ordinary dialogue in embedding space, that dimension features are partially independent, and that the optimal core–boundary configuration depends on the downstream model. LINGUAFORCE adopts the same dimension-as-anchor view but organizes it around the broader notion of agentive force of discourse: fifteen strategy types in four families measured on seven dimensions, with the existing coercion annotations serving as one of several validation targets.

B. Compliance-Gaining and Persuasion

Communication research has long studied how speakers gain compliance. Marwell and Schmitt [7] extracted sixteen compliance-gaining strategies (promising, threatening, moral appeal, debt, esteem, etc.), and Wiseman and Schenck-Hamlin [8] derived fourteen strategies via multidimensional scaling. Cialdini [9] organized the underlying mechanisms into six principles: reciprocity, commitment and consistency, social proof, authority, liking, and scarcity. In NLP, the closest operationalization is SemEval-2023 Task 3, which detects persuasion techniques in online news [3]; its techniques (loaded language, name-calling, appeal to fear, false dilemma, etc.) overlap with our families B–D but are defined for monological news text rather than interactive dialogue. Speech-act theory [11], [12] provides the directive class that anchors our taxonomy's transparent family. Our taxonomy is deliberately integrative: it maps these heterogeneous schemes onto one consistent label space.

C. Psychological Manipulation in Dialogues

Recent datasets target covert psychological manipulation. MentalManip annotates mental-manipulation techniques in dialogues [4]; MultiManip provides a multi-label dataset of manipulation types in real-world variety-show conversations and proposes a self-perception two-stage introspection method [5]; ReaMent collects real-world mental-manipulation dialogues [6]. These resources are valuable but narrow: they focus on covert or pathological strategies, omit benign and transparent control strategies (requests, advice, rational persuasion), and use task-specific taxonomies with incompatible label spaces. LINGUAFORCE differs by covering the full spectrum from transparent requests to covert abuse and by making all types comparable through one seven-dimension scale.

D. Toxicity, Condescension, and Implicit Abuse

Explicit toxicity detection is mature [1], and adjacent work targets condescension [2] and microaggressions [13]. These phenomena are part of our spectrum: verbal abuse and value shaming are two of our fifteen types, and toxicity is one of our seven dimensions. The key difference is scope: toxicity models treat hostility as the phenomenon, whereas we treat it as a channel through which agency is exerted, and we additionally require the discourse to target a listener's behavior or choice space.

E. Moral Foundations and Normative Language

Moral coercion operates by invoking moral norms, and our normative family (C) inherits that machinery. Moral foundations theory [14] identifies the dimensions—care, fairness, loyalty, authority, sanctity, and liberty—along which moral judgments vary, and computational work has shown that moral foundations are detectable in text and predictive of online discourse behavior [15]. We use these foundations to structure the annotation of D_3 (normative pressure) and the six types of family C: for example, C1 (moral appeal) typically loads on care and loyalty, whereas C4 (authority) loads on authority and liberty. This grounds the taxonomy in established psychological theory and provides features for the statistical analyses in Section IV. It also sharpens the boundary between C and the other families: normative pressure is the channel that most directly distinguishes moral coercion from, say, a threat (which loads on D_2 without invoking a moral foundation).

F. Agentive Force in LLM Safety Evaluation

Beyond detection, the agentive force of discourse is directly relevant to LLM safety. Current safety evaluations primarily measure explicit harms such as toxicity, bias, and refusal correctness; manipulative or coercive styles—guilt-heavy advice, authority-laden justifications, false-dilemma framings—are rarely scored even though they can steer users toward harmful decisions in persuasive domains such as health, finance, and politics. Benchmarks such as ToxicChat [1] and TalkDown [2] capture hostility and condescension in user–model interaction but do not measure the perlocutionary pressure of model outputs.

LINGUAFORCE's seven-dimension profile provides a continuous, interpretable measurement that can be applied to generated text: a model output can be scored for D_2 (option constraint) or D_3 (normative pressure) with the same rubric used for human dialogue. We therefore view the benchmark as a dual-purpose instrument: a detection benchmark for analyzing human conversation and a safety rubric for auditing model behavior.

G. Language and Agency in Linguistic Anthropology

Ahearn [16] introduced "language and agency" to linguistic anthropology, where agency denotes the socioculturally mediated capacity of a speaker to act. To avoid conflating the speaker-side notion with our listener-directed notion, we use the term agentive force of discourse (equivalently, perlocutionary pressure) throughout, and reserve "linguistic agency" for the anthropological sense.

III. Framework: A Unified Account of agentive force of discourse

A. Definition and Inclusion Criterion

We define the agentive force of discourse of an utterance or dialogue as its perlocutionary capacity to drive, guide, or restrict a listener's behavior and decision space. Formally, for an utterance u addressed to a listener L in context C, let $\mathcal{A}(u, C)$ be the set of actions that L considers acceptable after processing u, and let $\mathcal{A}^*(C)$ be the speaker-preferred subset. The agentive force of u is the pressure it exerts to shift $\mathcal{A}(u, C)$ toward $\mathcal{A}^*(C)$, along one or more of the seven channels defined below. Three properties are important. First, success is not required: perlocutionary force exists even when the listener resists. Second, force is compositional: an utterance can combine several channels (e.g., guilt plus social proof). Third, force is scalar: strategies differ in intensity, not only in kind.

The inclusion criterion follows directly: a discourse is agency-bearing if and only if it (i) targets a recognizable action or decision of the listener, (ii) employs at least one strategy in our taxonomy, and (iii) aims to drive or restrict behavior. Pure information statements and pure emotional venting that do not direct behavior are excluded and serve as negative examples. This criterion makes the negative class meaningful—a common weakness in manipulation benchmarks that sample only positive examples.

B. Three-Layer Architecture

Figure 1 shows the architecture. Layer 1 is a functional taxonomy: what the speaker is doing (fifteen strategy types in four families). Layer 2 is a measurement layer: how much pressure is applied along each of seven psychological dimensions. Layer 3 is the linguistic-realization layer: how the pressure is realized in surface form. The three layers are jointly annotated in LINGUAFORCE, which allows models to reason at the intended abstraction level and enables error analysis that is impossible with binary labels alone.

C. Strategy Taxonomy

Table I presents the taxonomy. Four families organize the strategies by their dominant pressure channel. Family A (transparent) contains low-pressure, intention-transparent strategies and provides the behavioral baseline of the spectrum. Family B (exchange) motivates behavior through anticipated benefit or harm. Family C (social-normative) exploits norms, roles, and group identity; it contains the moral-coercion strategies previously studied as a single task and expands them into a family of six types. Family D (covert) contains implicit or deceptive strategies that operate below the listener's awareness. The taxonomy is compatible with prior coercion annotations: obligation types (C_2) and constraint types (B_2, D_3) recover the obligation and constraint channels, moral appeal, shaming, and authority (C_1, C_3, C_4) cover the value-judgement channel, and verbal abuse (D_4) covers the toxicity channel. The remaining types extend the label space with strategies not covered by prior coercion work.

D. Universal Psychological Dimensions

Table II defines seven dimensions measured on a continuous 0–1 scale (with a companion four-level discrete label), at both turn and dialogue level. The inventory is designed to be complete with respect to the taxonomy: every strategy in Table I loads on at least one dimension, and the seven dimensions are intended as the minimal set that separates all fifteen types. It is also compatible with prior coercion annotations by construction: $D_2 \supseteq$ CS, $D_3 \supseteq$ OB $\cup$ VJ (normative pressure merges the two moral channels), and $D_6 \supseteq$ TX. Consequently, dialogues annotated with these channels can be embedded into the seven-dimensional space without loss of information, which is what makes the annotation-consistency checks in Section V well-defined.

Because the inventory is compatible with prior coercion annotations by construction (Table II), the reused dialogues double as a validation set for the new annotation scheme: dialogues labeled as morally coercive should receive high D_3 and low-to-moderate D_7. We also expect each strategy family to load primarily on a small set of dimensions—D_3 for family C, D_2/D_4 for family B, and D_5/D_6 for family D—which we examine through per-family dimension profiles and a clustering analysis (Section IV). These analyses serve as quality checks on the annotations and as descriptive evidence for the structure of the new benchmark.

E. Linguistic Realization Layer

Table III links each psychological dimension to its typical surface realizations. The layer serves two purposes. First, it makes the annotation operational: annotators use the cue checklist to justify dimension scores, which improves agreement. Second, it supports error analysis: when a model misclassifies a dialogue, we can trace whether the failure stems from missing cues, cue ambiguity, or

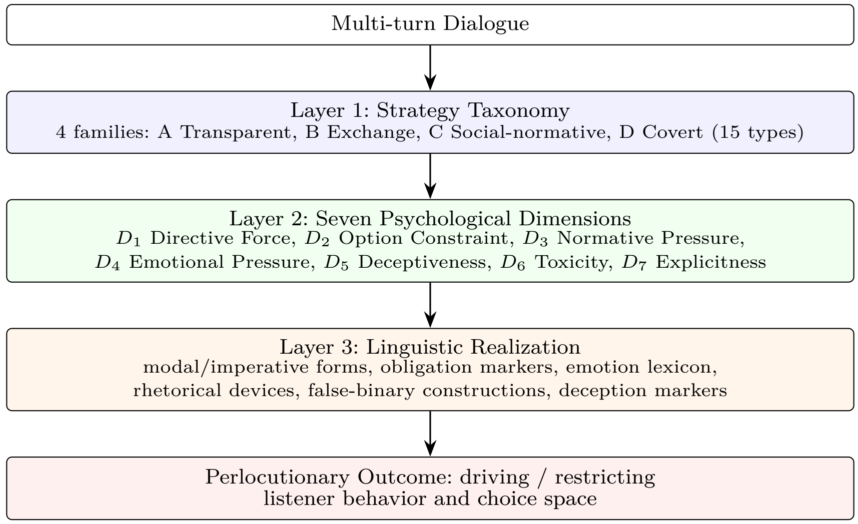

Fig. 1. Three-layer architecture of LINGUAFORCE. The strategy taxonomy defines what the speaker does, the psychological dimensions measure how much pressure is applied, and the linguistic-realization layer links both to surface cues. Moral coercion is a special case of this spectrum, located at high D_3 (normative pressure) with moderate D_1 and low-to-moderate D_7 (explicitness).

long-range context that surface forms cannot capture. The realization layer also connects LINGUAFORCE to explainability: dimension scores extracted by an upstream parser (Section V) are directly verbalizable as cue-based rationales.

F. A Worked Annotation Example

Table IV shows a two-turn dialogue and its gold annotation to make the protocol concrete. The dialogue mixes two strategies—moral appeal (C1) and obligation (C2)—and therefore receives a multi-label type annotation; its intensity is high, and the dimension profile is dominated by D_3 with secondary D_1. The example illustrates the two properties that make the task hard: the language is polite (no D_6) and the pressure is implicit (low D_7), exactly the pattern that defeats lexical toxicity classifiers.

G. Dimension Structure and Interdependence

The seven dimensions are not assumed independent; their correlation structure is itself an empirical object. We state the expected structure explicitly so it can be falsified. (i) D_1 (directive force) should correlate positively with D_2 (option constraint): the harder an utterance pushes, the more it tends to narrow options; the transparent family (A1) is the predicted exception, with high D_1 but low D_2. (ii) D_3 (normative pressure) and D_4 (emotional pressure) should co-occur in family C—guilt and moral appeal are usually emotionally loaded—but should separate in family B, where threats are affectively cold. (iii) D_5 (deceptiveness) should be roughly orthogonal to the other dimensions, since gaslighting can occur in both polite and hostile registers; a strong correlation with D_6 or D_7 would indicate an annotation bias rather than a property of the phenomenon. (iv) D_7 (explicitness) should correlate negatively with D_5: covert strategies are typically indirect. These predictions are tested in the statistical analysis plan (Section IV) via correlation matrices, factor analysis, and per-family conditional correlations; any deviation is reported and used to revise the guideline rather than silently discarded.

IV. Dataset Construction and Analysis

A. Design Principles

Four principles guide the construction. (i) Real-world multi-turn dialogues: the source corpus contains authentic multi-turn interaction rather than scripted dialogue, which lacks the fragmentation, topic drift, and pragmatic embedding of real conversation. (ii) Full-spectrum coverage: the corpus spans coercive and non-coercive dialogue across the full 0–5 intensity range, so benign controls are always present and degenerate models that classify any request as manipulation are prevented. (iii) Comparable measurement: every instance receives the same seven-dimension annotation, so all dialogues are comparable on one scale. (iv) Hard negatives: polite coercive utterances, sarcastic requests, and superficially kind manipulation are naturally present in the corpus and stress-test lexical shortcuts.

B. Sources and Sampling

LINGUAFORCE is built on a real-world corpus of everyday multi-turn dialogues collected by prior work [10]. We take the corpus's dialogues as raw text, retain their existing binary coercion and 0–5 intensity labels for validation, and re-annotate every dialogue under our unified scheme with the seven dimensions and a dialogue-level intensity. The full release comprises the 3,432-dialogue training partition; a disjoint 634-dialogue held-out partition is released as a first consistency-check set (Section IV-G), and its statistics are reported alongside the full set. The dialogues are already de-identified and

TABLE I
Strategy taxonomy of LINGUAFORCE: four families and fifteen types. Prior-channel anchors show the correspondence to previously studied coercion dimensions (OB/CS/VJ/TX).

Family	Type	Description and example	Prior	Role
A. Transparent	A1 Request/Advice	Low-pressure, intention-transparent proposal. "Could you revise this report by Friday?"	—	benign control
	A2 Directive/Order	Direct instruction with expected compliance. "Send the report before you leave."	—	transparent pressure
	A3 Rational Persuasion	Argument, evidence, expert opinion. "The data show 30% lower cost; adopt it."	—	reasoned pressure
B. Exchange	B1 Promise/Inducement	Benefit-driven motivation. "Help me this once and I will return the favor."	—	incentive pressure
	B2 Threat/Warning	Harm-driven motivation. "If you do not sign, you will face consequences."	CS (explicit)	explicit constraint
C. Social-normative	C1 Moral Appeal/Guilt	Appeals to conscience or moral duty. "A dutiful son would never let his mother suffer."	VJ+OB	normative core
	C2 Obligation/Duty	Invokes role responsibilities. "You are the team leader; it is your duty."	OB	duty pressure
	C3 Value Judgement/Shaming	Moral labeling of the target. "You are so selfish; you never think of others."	VJ	moral attack
	C4 Authority	Appeals to status, rank, seniority. "The boss decided; why keep arguing?"	—	hierarchical pressure
	C5 Conformity/Social Proof	Appeals to group norms. "Everyone is doing it; why are you special?"	—	peer pressure
	C6 Reciprocity/Debt	Invokes past favors. "I helped you last time; do not refuse this."	—	debt pressure
D. Covert	D1 Emotional Manipulation	Feigned weakness, pity, or fear to steer behavior. "If you leave too, I do not know what to do."	—	affective leverage
	D2 Deception/Gaslighting	Distorts facts or memory. "I never said that; you misremember."	—	epistemic attack
	D3 Logical Trap/False Dilemma	Shrinks the option space. "Either you follow me or we are done."	CS	option elimination
	D4 Verbal Abuse/Toxicity	Explicit hostility and disparagement. "You are useless; you cannot even do this."	TX	hostile pressure

TABLE II
Seven psychological dimensions and their relation to previously studied coercion channels (OB/CS/VJ/TX).

Dim.	Definition	Prior
D_1 Directive Force	Strength with which the utterance drives a target action	generalization (new)
D_2 Option Constraint	Degree to which the listener's options are narrowed	$\supseteq$ CS
D_3 Normative Pressure	Pressure via duty, morality, norms, roles	$\supseteq$ OB+VJ
D_4 Emotional Pressure	Use of guilt, shame, fear, pity as leverage	new
D_5 Deceptiveness	Degree of factual or epistemic distortion	new
D_6 Toxicity	Explicit hostility and abuse	$\supseteq$ TX
D_7 Explicitness	Transparency vs. indirectness of the pressure	new

cover everyday scenarios, so the negative class (non-coercive dialogue) and the intensity tail arise naturally; no additional collection is required for this release.

TABLE III
Illustrative surface cues per dimension. Cues are indicative, not sufficient: a cue only counts when the discourse targets the listener's behavior.

Dim.	Typical cues
D_1	imperative and hortative forms; modal "must/should/have to"; performative requests
D_2	"either/or", ultimatums, no-alternative framings, threats of consequence
D_3	duty and role nouns, moral adjectives, normative "ought"; appeal to what "a good X" would do
D_4	emotion lexicon (guilt, shame, fear, pity), emotional blackmail, self-deprecation
D_5	fact distortions, memory disputes, gaslighting patterns, manufactured consensus
D_6	insults, disparagement, profanity, dehumanizing labels
D_7	direct address and explicit request vs. hints, presuppositions, innuendo

C. Annotation Protocol

Annotation follows a three-level decision tree: (1) Does the discourse target a listener action or decision? If no, it is a negative example. (2) Which strategies from Table I are present? (multi-label; a dialogue may combine, e.g., moral appeal and social proof). (3) What are the intensity

TABLE IV
A worked annotation example. S = speaker turn; L = listener turn. Type, intensity, and the seven dimension scores are annotated per turn and aggregated at dialogue level.

Turn	Utterance	Type (multi-label)	Dimension profile
S1	"A dutiful son would never let his mother live alone; you are not that kind of person, are you?"	C1 Moral Appeal; C3 Value Judgement	D_3=0.9, D_1=0.7, D_4=0.6, D_7=0.4
L1	"I know, but my job requires me to move..."	— (defensive)	D_1=0.2
S2	"Your mother sacrificed everything for you; it is your turn to repay her."	C2 Obligation; C6 Reciprocity	D_3=0.9, D_2=0.5, D_1=0.8
Dialogue-level			Intensity 4/5; T1 positive; C1+C2+C3+C6

TABLE V
Released LINGUAFORCE sets.

Set	Dialogues	Coercive	Non-coercive
Full release	3,432	1,862 (54.2%)	1,570 (45.8%)
First-release held-out	634	332 (52.4%)	302 (47.6%)

and the seven dimension scores? Intensity is ordinal 0–5; each dimension receives a continuous 0–1 score and a four-level discrete label (None/Low/Moderate/High). In this release the annotation is produced by an instruction-tuned LLM (Section V) that follows the decision tree and emits a structured JSON profile for each dialogue at temperature 0 for reproducibility. Reliability is validated through the consistency checks against the existing labels in Section IV-G and through a two-annotator agreement study on a stratified subsample (Section IV-E).

D. Annotation Decision Tree and Negative Cases

To keep the negative class principled, annotation follows the decision tree in Figure 2. The first question filters out the majority of ordinary conversation; the second assigns strategy labels; the third assigns intensity and the seven dimension scores. Negative examples are annotated at the same granularity as positives so that the dimension space is defined over all dialogues, not only manipulative ones—this is what allows the dimension space to represent the full spectrum, including non-manipulative dialogue. Typical negative cases include pure information exchange ("The meeting is at three"), emotional venting with no behavioral target ("I am completely exhausted"), and questions that seek information rather than action ("What time is the deadline?"). Borderline cases, such as advice that is also a subtle request, are resolved by the multi-label rule: assign every applicable type, and let the dimension scores capture the dominant channel.

E. Quality Control and Consistency Validation

Because annotation is LLM-produced rather than crowd-sourced, quality control combines measurable consistency with the source corpus's gold labels with a manual re-annotation check. We report, for every dimension, its Spearman correlation with the binary coercion label and the 0–5 intensity label, and the AUC for separating coercive from non-coercive dialogues (Tables VII and VIII); a dimension that fails to track the gold labels is flagged as an unreliable anchor. We additionally check that the predicted intensity is monotonic in the gold label (Figures 5 and 9) and that the inter-dimension correlation matrix shows no pathological coupling (Figures 6 and 10). To probe reliability beyond label consistency, the first author manually annotated a stratified subsample of 120 dialogues (balanced on the gold binary label and spanning all intensity levels), following the same annotation rubric as the automatic pipeline. Agreement with the automatic profiles is high on overall agency strength (Spearman ρ=0.79), and the primary pressure mechanism matches the automatic profile's top two dimensions in 81.2% of the 85 mechanism-bearing cases (excluding the explicitness meta-dimension D_7, which marks directness rather than a specific pressure channel). These results indicate that the seven-dimension profiles are reproducible against independent human judgment.

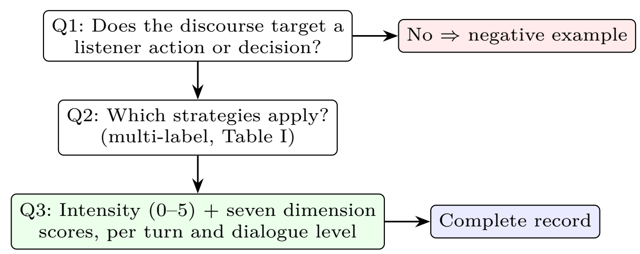

Fig. 2. Three-question annotation decision tree.

To further probe reliability beyond the author's spot-check, we ran a two-annotator agreement study on a stratified subsample of 150 dialogues (balanced on the binary label), annotated independently under the same rubric. Table VI reports inter-annotator agreement on the discrete labels: the core intensity label reaches substantial agreement (quadratic-weighted κ=0.73), as do normative pressure (0.78), toxicity (0.76), and option constraint (0.66); binary presence and explicitness reach moderate agreement (0.49), and deceptiveness is the least reliable dimension (0.39), consistent with its subjective, low-

TABLE VI

Inter-annotator agreement on a stratified subsample (n=150; two independent annotators, 149 fully labeled pairs). Binary uses unweighted Cohen's κ; intensity and dimensions use quadratic-weighted κ.

Variable	Cohen's κ	Agreement level
binary (pressure present)	0.49	moderate
intensity (0–5)	0.73	substantial
D_1 directive force	0.55	moderate
D_2 option constraint	0.66	substantial
D_3 normative pressure	0.78	substantial
D_4 emotional pressure	0.60	substantial
D_5 deceptiveness	0.39	fair
D_6 toxicity	0.76	substantial
D_7 explicitness	0.49	moderate

frequency nature. The continuous scores inherit the same ordering as the discrete labels, so the reported agreement carries over to the released scores.

F. Statistical Analyses

We report: (i) per-dimension means and consistency with the gold labels (Table VIII); (ii) the dimension correlation matrix, testing whether D_3–D_4 co-occur while D_5 stays orthogonal (Figures 6 and 10); (iii) the separation of coercive and non-coercive dialogues in the seven-dimensional space (Figures 4 and 8); and (iv) the monotonicity of predicted intensity in the gold 0–5 label (Figures 5 and 9). These analyses also serve as sanity checks against spurious lexical correlations, a failure mode identified in prior coercion datasets. Type-level analyses on the type-labeled split are reported in Section V-H.

G. First Release and Consistency with Existing Labels

To validate the framework before scaling annotation, we release a first consistency-check set: the 634-dialogue held-out split of the underlying corpus, re-annotated under the unified seven-dimension scheme. The set contains 332 coercive (52.4%) and 302 non-coercive dialogues, and gold intensity is distributed as 170/54/78/64/107/161 over levels 0–5, so hard negatives and the high-intensity tail are both represented.

Table VII and Figures 3–6 report the consistency of the extracted profiles with the existing labels. All seven dimensions correlate with the binary coercion label and with the 0–5 intensity label; the strongest channels are option constraint (D_2), directive force (D_1), and emotional pressure (D_4), followed by explicitness (D_7) and normative pressure (D_3); this is consistent with constraint being a core channel of moral coercion [10]. Each dimension separates coercive from non-coercive dialogues with AUC between 0.72 and 0.83. The predicted intensity rises monotonically with the gold label (Spearman ρ=0.66), and the inferred binary decision reaches AUC 0.83 and F1 0.80 at the best threshold, confirming that the unified dimension space reproduces—and can in principle subsume—the prior corpus's own measurement on the

TABLE VII

Consistency of the seven dimensions with the existing labels on the first release (n=634). ρ: Spearman correlation; AUC: area under the ROC curve for separating coercive from non-coercive dialogues.

Dimension	ρ (binary)	ρ (0–5)	AUC
D_1 Directive force	0.54	0.60	0.81
D_2 Option constraint	0.57	0.63	0.83
D_3 Normative pressure	0.49	0.63	0.78
D_4 Emotional pressure	0.54	0.60	0.81
D_5 Deceptiveness	0.49	0.55	0.77
D_6 Toxicity	0.42	0.35	0.72
D_7 Explicitness	0.53	0.53	0.79
Predicted intensity	–	0.66	0.83

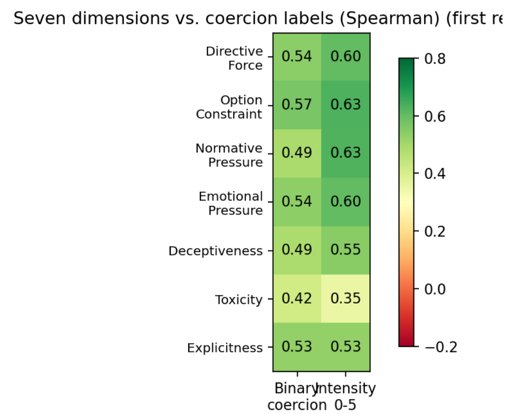

Fig. 3. Seven dimensions versus the existing binary and 0–5 intensity labels (Spearman correlation).

shared dialogues. A direct zero-shot judge that bypasses the profile achieves AUC 0.78 and best F1 0.76 on the same dialogues (Section V-H), confirming that the added benefit of the seven-dimension protocol is modest on this already-annotated split but consistent on the harder binary decision. Section IV-H repeats the same analysis on the full 3,432-dialogue corpus.

H. Full-Scale Annotation and Dimension Profiles

Having validated the protocol on the held-out split, we scale the seven-dimension annotation to the complete corpus: the full LINGUAFORCE release contains all 3,432 dialogues, of which 1,862 (54.2%) are coercive and 1,570 (45.8%) are non-coercive. Gold intensity is distributed as 986/267/317/337/621/904 over levels 0–5; hard negatives (level 0, 28.7%) and the high-intensity tail (levels 4–5, 44.4%) together dominate the corpus, so the benchmark remains challenging for both the binary and the ordinal decision.

Table VIII summarizes the seven dimensions on the full set. Option constraint (D_2) and emotional pressure (D_4) remain the strongest channels (AUC 0.815 and 0.829), followed by directive force (D_1, 0.792) and explicitness (D_7, 0.788). Deceptiveness (D_5, 0.767) and toxicity (D_6, 0.710) are weaker, and toxicity in particular shows the

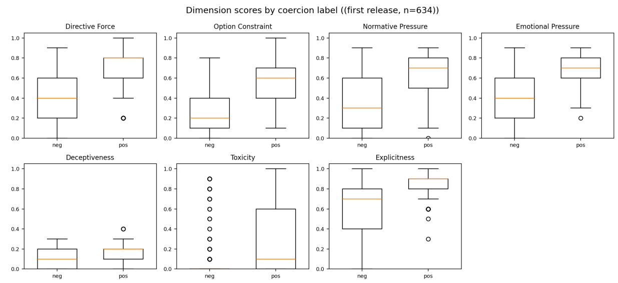

Fig. 4. Dimension scores for coercive versus non-coercive dialogues (n=634).

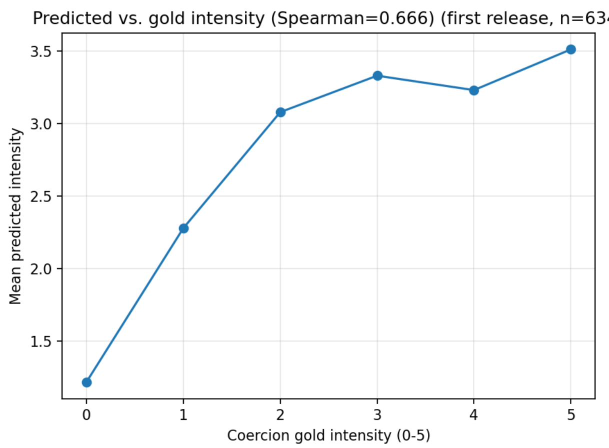

Fig. 5. Predicted intensity as a function of the gold 0–5 intensity label; points are within-level means.

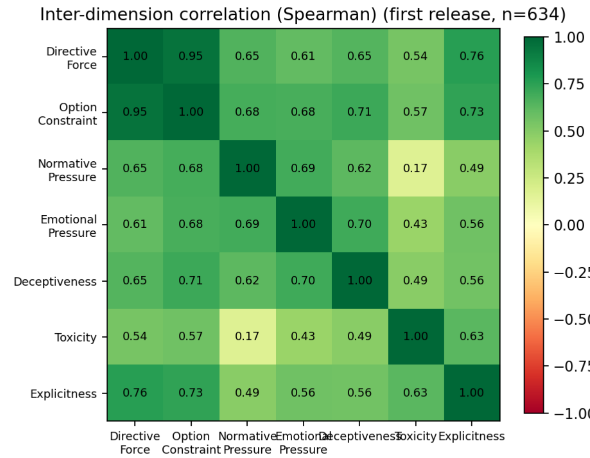

Fig. 6. Inter-dimension correlation matrix (Spearman).

TABLE VIII
Seven-dimension profiles on the full release (n=3,432). Mean score, Spearman correlation ρ with the binary label and the 0–5 intensity label, and AUC for separating coercive from non-coercive dialogues.

Dimension	Mean	ρ (binary)	ρ (0–5)	AUC
D_1 Directive force	0.53	0.51	0.54	0.79
D_2 Option constraint	0.36	0.55	0.58	0.82
D_3 Normative pressure	0.49	0.45	0.55	0.76
D_4 Emotional pressure	0.53	0.57	0.62	0.83
D_5 Deceptiveness	0.12	0.48	0.50	0.77
D_6 Toxicity	0.22	0.40	0.31	0.71
D_7 Explicitness	0.73	0.51	0.49	0.79
Predicted intensity	–	0.63	0.65	0.85

lowest correlation with the coercion labels (ρ=0.40 binary, 0.31 intensity): abusive language overlaps with but does not define the coercive core. Since no single channel reaches AUC 0.83 on its own, the coercion decision is not reducible to one surface feature, which motivates the multi-dimensional representation. The predicted overall intensity stays monotonic in the gold 0–5 label across the full corpus (Spearman ρ=0.65; AUC 0.85).

I. Ethics and Privacy

Data provenance and consent. The dialogues inherit the source corpus's de-identification process: the corpus was already anonymized by prior work [10], and this release adds no new personal information. Collection of the underlying corpus complied with the platform terms of service in force at the time; this paper performs re-annotation of existing text and performs no new data collection, so no additional informed consent is required beyond what the source project obtained.

IRB and human subjects. The study involves no human participants: all annotations are produced by a frozen instruction-tuned LLM following a public decision tree, and no behavioral or demographic data are collected. A human spot-check of 120 dialogues was performed by the first author on already-anonymized text and is not research on human subjects; an IRB review is therefore not required under our institutional policy.

Annotation welfare. Although the bulk of annotation is automated, the spot-check and any future human annotation expose annotators to potentially distressing content (e.g., verbal abuse, guilt-tripping). Annotators receive written guidelines, may skip any instance without penalty, and take breaks; the released guidelines accompany the code repository.

Dual-use and abuse risk. The taxonomy and dimensions could, in principle, be used to craft more effective manipulative language. We release the dataset and code for research purposes only, emphasize detection and analysis rather than generation, and provide no generation recipes. The dataset is not intended for automated moderation of real users without human review, for surveillance, or for consequential decisions about individuals.

Culture and language scope. The current release is English and dominated by everyday scenarios. The seven-dimension scheme is grounded in pragmatics and persuasion research, but its cross-lingual and cross-cultural validity is an open question; we caution against treating the dimensions as culture-universal without further study.

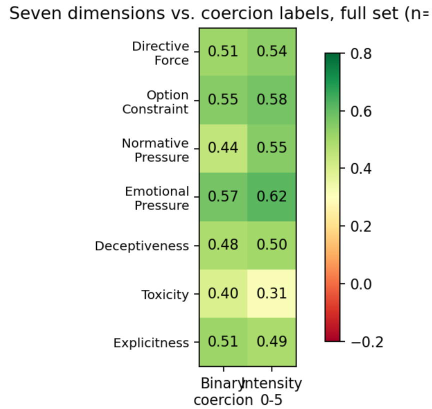

Fig. 7. Seven dimensions versus the binary and 0–5 intensity labels on the full release (n=3,432; Spearman correlation).

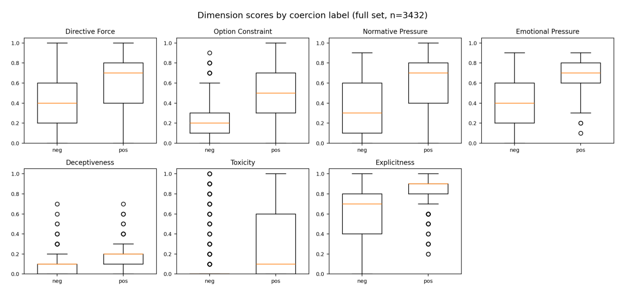

Fig. 8. Dimension scores for coercive versus non-coercive dialogues on the full release (n=3,432).

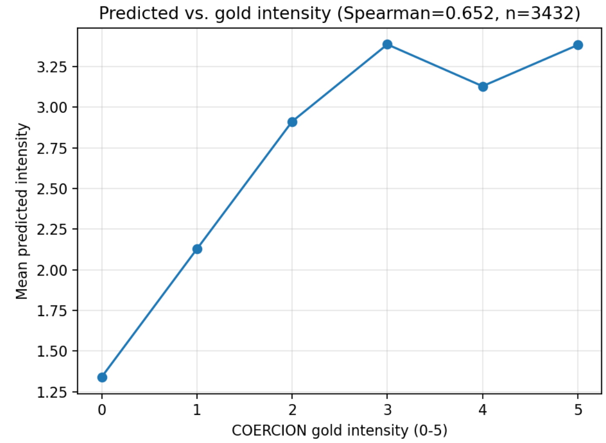

Fig. 9. Predicted intensity as a function of the gold 0–5 intensity label on the full release; points are within-level means.

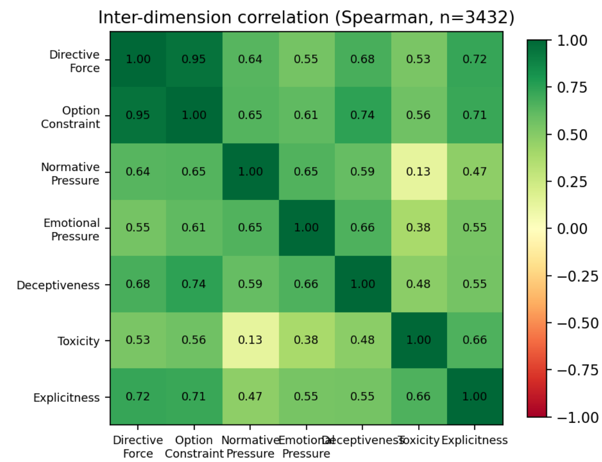

Fig. 10. Inter-dimension correlation matrix on the full release (Spearman).

Data statement and maintenance. A full datasheet [17] accompanies the release (datasheet.md), documenting composition, collection, preprocessing, intended uses, and known limitations. The dataset is released under a research-only license; corrections are tracked as versioned releases with a changelog.

V. Experimental Design and Evaluation Protocol

A. Tasks

LINGUAFORCE defines three tasks at increasing granularity. T1 (manipulation detection) is a binary classification of whether a dialogue exerts agentive force beyond the transparent baseline. T2 (type recognition) is a fifteen-way multi-label classification of the strategies in Table I; type labels are released for both the held-out split and the full release. T3 (intensity prediction) is ordinal regression of the 0–5 intensity, evaluated with quadratic weighted kappa (QWK) and accuracy-within-±1.

B. Models and the Two-Stage Pipeline

Figure 11 shows the two-stage dimension-aware pipeline. An upstream LLM parses a dialogue into an agentive profile, the seven-dimension vector at both turn and dialogue level; a downstream model consumes the profile as structural anchors and predicts T1/T2/T3. This two-stage design decouples dimension extraction from classification. We compare three aggregation modes—turn-level only (T), global only (G), and turn+global (B)—since aggregation granularity is known to affect dimension-based classification.

The pipeline is evaluated in a zero-shot setting: a frozen instruction-tuned LLM (DeepSeek) acts as the upstream dimension parser, and a direct judge that skips the profile serves as the control. Prompt variants include zero-shot, chain-of-thought [18], and self-reflection, following the self-perception approach of MultiManip [5]. Table IX summarizes the configuration.

C. Experimental Setup

All evaluation in this release is zero-shot: a frozen instruction-tuned LLM (DeepSeek) parses each dialogue

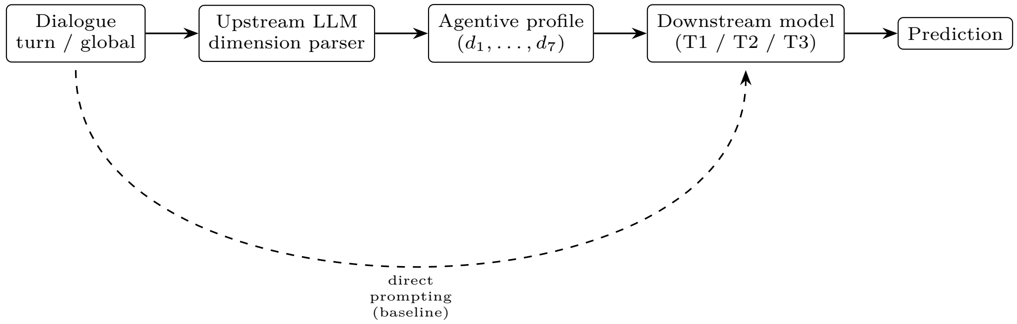

Fig. 11. Two-stage dimension-aware pipeline. The dashed path is the direct prompting baseline used to isolate the effect of dimension anchoring.

TABLE IX
Model configuration for the zero-shot evaluation in this release.

Model	Role	Tuning
DeepSeek	upstream dimension parser	frozen
DeepSeek	direct zero-shot judge	frozen

into the seven-dimension profile and a dialogue-level intensity at temperature 0 for reproducibility, and T1 and T3 are read off the profile directly, with no fine-tuning and no in-domain training data. We compare against a direct zero-shot judge that skips the profile (dashed path in Fig. 11). Metrics are binary AUC and best F1 over the intensity threshold for T1, and Spearman, Pearson, QWK, and accuracy-within-± 1 for T3. Cross-domain transfer uses the same frozen parser on four public corpora (Section V). The parser runs through a public API at low cost; no GPU cluster is required.

D. Research Questions

RQ1 (annotation consistency): Do the seven dimensions align with existing coercion annotations on the shared dialogues? Protocol: on the 634-dialogue first release and the full 3,432-dialogue release (Sections IV-G–IV-H), report per-dimension correlation with the existing binary coercion label and the 0–5 intensity label, together with the AUC of each dimension for separating coercive from non-coercive dialogues. High consistency shows that the new annotation scheme reliably covers previously studied phenomena.

RQ2 (dimension structure): Do the seven dimensions carry distinct, complementary signal? Protocol: report mean dimension scores and per-dimension AUC for separating coercive from non-coercive dialogues, and examine the inter-dimension correlation matrix for expected couplings (e.g., D_3–D_4) and unexpected ones (e.g., D_5–D_6). Per-family and per-type analyses on the type-labeled split are reported in Section V-H.

RQ3 (zero-shot generalization): Can dimension conditioning enable detection of unseen manipulation types? The primary evidence is cross-domain zero-shot transfer to corpora whose manipulation types are never seen during training (Section V): AUROC 0.742 on MentalManip, 0.706 on MultiManip, and 0.729 on TalkDown. We additionally ran the within-taxonomy leave-one-type-out protocol on the type-labeled held-out split: a model trained only on the other fourteen types cannot separate a held-out type as a novel cluster in the seven-dimension space (mean novelty AUROC 0.54, near chance), consistent with the framework's design in which types are graded on shared continuous dimensions rather than forming disjoint clusters. Crucially, the dimension-conditioned pipeline still flags dialogues of unseen types as manipulative: predicted intensity separates dialogues containing any strategy from benign dialogues with AUROC 0.963. The framework therefore transfers to unseen manipulation types as a pressure signal, while assigning the specific type label remains an open challenge.

RQ4 (engineering conclusions): How do aggregation mode and prompt design interact with the dimension-conditioned pipeline? We run two slices of the grid with the same linear readout as Table XII (trained on the full release, evaluated on the held-out split). Aggregation (Table X): global-only (G), turn-mean (T), and their concatenation (B). Turn-mean pooling collapses toward chance (binary AUC 0.525) because manipulative pressure is concentrated in a minority of turns and averaging dilutes it; the global profile (G) recovers the signal, and concatenation (B) performs best (binary AUC 0.859, intensity Spearman 0.679, QWK 0.680). Prompt variants (Table XI): zero-shot, chain-of-thought, and self-reflection yield comparable downstream performance (AUC 0.83–0.84), showing the pipeline is robust to prompt design; CoT marginally helps T1 and self-reflection T3. The remaining grid dimensions (parser-model choice and ICC fidelity) are left to future work.

E. Cross-Domain Evaluation

To test transfer and distinctiveness, we evaluate the zero-shot pipeline on existing resources with compatible labels: MentalManip [4] and MultiManip [5] for psychological manipulation, TalkDown [2] for patronizing language,

TABLE X
Aggregation ablation of the dimension-conditioned pipeline (linear readout; train full release, held-out split). T mean-pools the per-turn profiles over turns; B concatenates the turn-mean and global profiles.

Aggregation	Binary AUC	Intensity ρ	QWK
G (global only)	0.855	0.674	0.667
T (turn mean)	0.525	0.040	0.066
B (turn+global)	0.859	0.679	0.680

TABLE XI
Prompt ablation of the upstream dimension parser (held-out split, n=634). All variants use the same downstream linear readout.

Prompt variant	T1 AUC	T3 Spearman	T3 QWK
Zero-shot	0.832	0.666	0.623
Chain-of-thought	0.839	0.639	0.613
Self-reflection	0.838	0.670	0.630

and ToxicChat [1] for generic toxicity. We report each corpus's standard metric under its own decision rule and the threshold-free AUROC, which is directly comparable across label schemes. ReaMent [6] and SemEval-2023 Task 3 [3] are left as future work. We use the confusion across these corpora to identify which constructs are shared with agency and which are distinct.

F. Dimension Contribution and Anchoring

In this release, the contribution of each dimension is measured by its consistency with the gold labels: per-dimension Spearman correlation and AUC against the existing binary and 0–5 intensity labels (Tables VII and VIII). Dimensions with weak consistency (e.g., toxicity, D_6) are weaker anchors for the coercion decision, while strong ones (e.g., option constraint, D_2; emotional pressure, D_4) carry most of the signal. To quantify each dimension's marginal contribution, we fit a lightweight linear readout over the seven dimension scores—logistic regression for the binary decision and linear regression for the 0–5 intensity—trained on the full release and evaluated on the held-out split. The linear readout reaches binary AUC 0.855 and intensity Spearman 0.674 on the held-out set, slightly above the LLM's dimension-conditioned readout on the same split (AUC 0.831, Spearman 0.665); the annotated profiles alone therefore carry the decision-relevant signal and support cheap downstream models without an LLM. Table XII reports leave-one-dimension-out results: removing toxicity (D_6) degrades detection most (AUC drop of 0.008), removing normative pressure (D_3) degrades intensity most (Spearman drop of 0.021), and option constraint (D_2) also contributes. Notably, D_6 has the weakest univariate correlation yet its removal hurts detection most, indicating it supplies non-redundant information (e.g., separating hostile-but-non-coercive dialogues). Emotional pressure (D_4) is largely redundant with D_3 in the linear readout, and the remaining dimensions

TABLE XII
Leave-one-dimension-out ablation with a linear readout over the seven dimension scores (held-out split, n=634). ρ: Spearman correlation of predicted intensity with the gold 0–5 label. The LLM readout is the dimension-conditioned zero-shot reference.

Readout	Binary AUC	Intensity ρ
7-dim linear readout	0.855	0.674
$-D_1$ (directive force)	0.854	0.674
$-D_2$ (option constraint)	0.851	0.668
$-D_3$ (normative pressure)	0.851	0.653
$-D_4$ (emotional pressure)	0.860	0.693
$-D_5$ (deceptiveness)	0.855	0.674
$-D_6$ (toxicity)	0.847	0.669
$-D_7$ (explicitness)	0.853	0.674
LLM readout (reference)	0.831	0.665

have small marginal effects, consistent with the inter-dimension correlations of Figure 6.

G. Dimension-Space Structure

To verify that the seven dimensions carve distinct, interpretable regions for the four strategy families, Figure 12 embeds the 3,432 dialogue profiles in two dimensions (t-SNE), colored by the dominant family, together with the 226 dialogues for which no strategy was detected (benign). Families occupy distinguishable regions, and benign dialogues cluster at the low-pressure side. A logistic classifier over the seven raw dimensions separates the five groups with 0.640 accuracy (chance 0.2) in five-fold cross-validation, and each dimension alone separates benign from manipulative dialogues with AUC 0.73–0.94. Figure 13 reports each family's mean dimension profile: the social-normative family peaks on normative and emotional pressure (D_3, D_4), the exchange family on option constraint and toxicity (D_2, D_6), the covert family on emotional pressure (D_4), and benign dialogues are uniformly low across all dimensions.

H. Results on the Full Release

Table XIV and Table XV report the zero-shot results of the dimension-conditioned pipeline on the full 3,432-dialogue release: a dialogue is decoded into the seven-dimension profile (Section V), the profile is mapped to a 0–5 intensity, and T1 and T3 are read off directly. As a control, Table XIV also reports a direct zero-shot judge that skips the dimension profile (dashed path in Fig. 11) on the first-release split. The dimension-conditioned readout reaches binary AUC 0.852 and best F1 0.817 on the full release, exceeding the direct judge (AUC 0.780, best F1 0.755) and the first-release estimate (AUC 0.832, F1 0.795); the fine-grained profile therefore carries information that a single holistic judgment misses, and the estimate is stable as the corpus grows. On the 0–5 intensity scale the readout is well calibrated (QWK 0.595, Spearman 0.651, accuracy within ± 1 of 0.692 on the full release), with only a modest drop from the first-release values as hard negatives and the intensity tail are added.

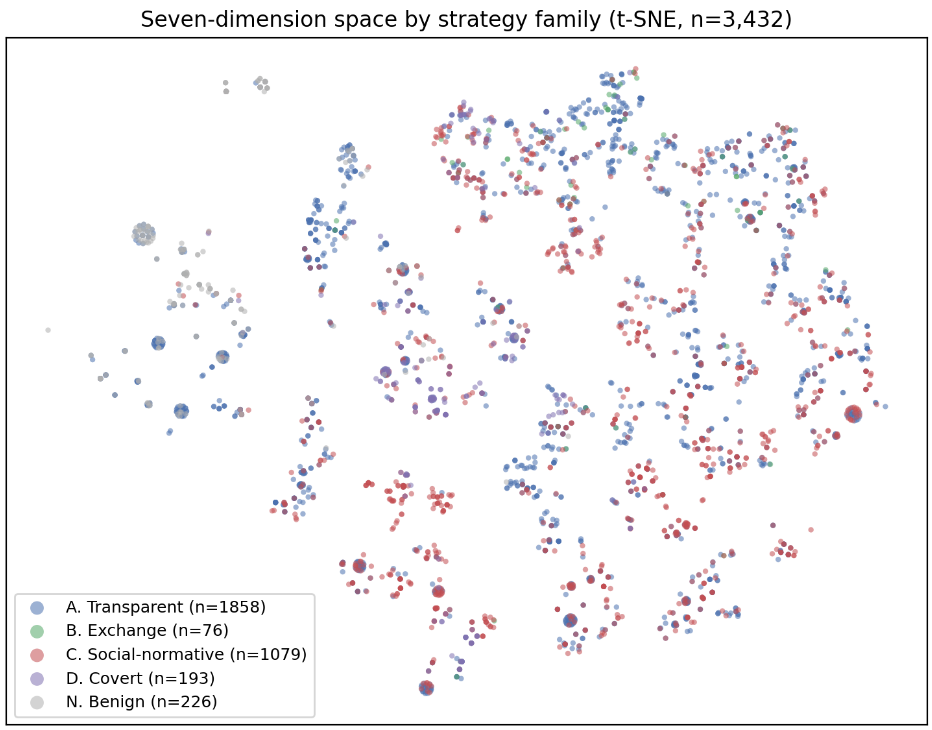

Fig. 12. t-SNE embedding of the seven-dimension profiles (n=3,432), colored by dominant strategy family; the 226 no-strategy dialogues form the benign group.

The cross-domain transfer results (Table XVI) support the interpretability and discriminative-validity claims. Under a threshold-free ranking, the zero-shot judge transfers moderately to the psychological- manipulation corpora (AUROC 0.742 on MentalManip, 0.706 on MultiManip) and to patronizing language (AUROC 0.729 on TalkDown), while transfer to generic toxicity is weak (AUROC 0.587 on ToxicChat). This pattern is consistent with the design: manipulation and condescension exert listener-directed pressure and are therefore partially captured by the agency dimensions, whereas generic toxicity is largely distinct from agency. Fixed-threshold macro-F1/F1 values are much lower (e.g., 0.058 on TalkDown) because the intensity scale is calibrated on our annotation scheme and the judge rarely assigns intensity $\geq$3 to dialogues annotated under other corpora's schemes: ranking (AUROC) transfers, absolute calibration does not. Target-calibrated thresholds recover macro-F1 around 0.68–0.74 on the manipulation corpora. T2 type recognition is evaluated with a lightweight one-vs-rest logistic readout over the seven dimension scores, contrasted with a supervised baseline that fine-tunes RoBERTa-base end-to-end on the dialogue text (Table XIII). We report two protocols for the readout. First, it is trained on the full 3,432-dialogue release and evaluated on the type-labeled held-out split: family-level recognition reaches micro-F1 0.827 and macro-F1 0.772, and type-level recognition reaches micro-F1 0.668 and macro-F1 0.480. Five-fold cross-validation over the full release yields similar estimates (family micro-F1 0.820 and macro-F1 0.752; type micro-F1 0.669 and macro-F1 0.477), so the results are stable as the corpus grows from 634 to 3,432 dialogues. The fine-tuned RoBERTa baseline gives consistent cross-validation estimates as well (type macro-F1 0.540, micro-F1 0.734; family macro-F1 0.767, micro-F1 0.815), confirming that the held-out results are not sensitive to the train/test split. The fine-tuned RoBERTa-base is clearly stronger at type level (macro-F1 0.544, micro-F1 0.731 on the held-out split), confirming that the type labels are learnable by a supervised model; at family level it matches the readout (macro-F1 0.780, micro-F1 0.820), so the dimension readout remains a competitive lightweight and interpretable baseline. Type-level macro-F1 is depressed by three rare types (rational persuasion A_3, authority C_4, deception/gaslighting D_2), which appear in fewer than thirty of the 3,432 dialogues; family-level results are the more reliable target for this corpus, and the full type labels are released for downstream use.

I. Limitations and Threats to Validity

Several limitations should be acknowledged. First, the seven dimensions are a design choice; a finer or coarser decomposition is possible, and the dimension structure may be sensitive to it. We mitigate this by reporting consistency against the existing coercion labels at both the first-release and full-scale level. Second, cross-platform variation in norms and politeness conventions may limit

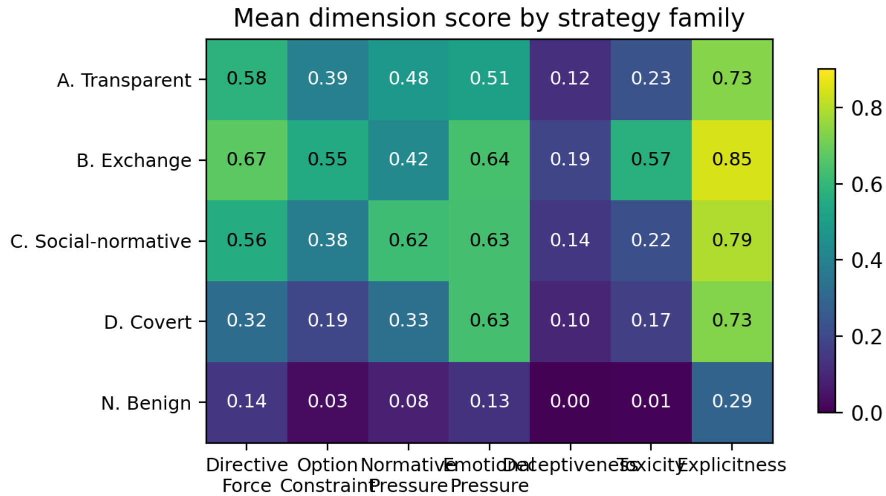

Fig. 13. Mean dimension score per strategy family. Social-normative and covert families concentrate pressure on normative/emotional dimensions; benign dialogues are uniformly low.

TABLE XIII

T2 type recognition: a one-vs-rest logistic readout over the seven dimension scores versus fine-tuned RoBERTa-base. "Held-out": trained on the full release (n=3,432), evaluated on the type-labeled held-out split (n=634); "5-fold CV": cross-validation over the full release. F1 is macro/micro at the best threshold.

		Linear readout		RoBERTa-base
Granularity	Protocol	Macro-F1	Micro-F1	Macro/Micro-F1
15 types	Held-out	0.480	0.668	0.544 / 0.731
15 types	5-fold CV	0.477	0.669	0.540 / 0.734
4 families	Held-out	0.772	0.827	0.780 / 0.820
4 families	5-fold CV	0.752	0.820	0.767 / 0.815

TABLE XIV

T1 detection on the full release (n=3,432): the dimension-conditioned readout, versus the same readout on the first-release split (n=634) and a direct zero-shot judge that skips the profile (dashed path in Fig. 11). "Best F1" is the maximum F1 over the intensity threshold.

Metric	Full (n=3,432)	First release	Direct zero-shot
Binary AUC	0.852	0.831	0.780
Best F1 (threshold)	0.817 (3)	0.795 (3)	0.755 (3)
Accuracy @ best F1	0.789	0.760	0.711
Best per-dimension AUC (D_2)	0.829	0.825	—

the transferability of both annotations and models; the cross-domain evaluation and a planned cross-lingual extension bound this risk. Third, annotation of covert strategies such as gaslighting is inherently subjective; because the current annotation is LLM-produced, we rely on consistency with the source corpus's gold labels as the reliability check, and we flag that a human agreement study is still needed. Fourth, the dataset inherits the sampling biases of the source corpus and of the platforms it

TABLE XV

T3 intensity prediction from the dimension-conditioned pipeline on the full release (n=3,432), with the first-release (n=634) values for comparison.

Metric	Full	First release
Spearman ho	0.651	0.665
Pearson r	0.677	0.691
Quadratic weighted kappa	0.595	0.623
Accuracy within ± 1	0.692	0.722
Exact match	0.216	0.219

TABLE XVI

Cross-domain zero-shot transfer of the seven-dimension pipeline. "Standard metric" is each corpus's own metric under the source-task decision rule (intensity ≥ 3); "AUROC" is threshold-free and directly comparable. Sampled subsets: MentalManip n=300, MultiManip n=220, TalkDown n=652, ToxicChat n=300.

Corpus	Standard metric	Zero-shot	AUROC
MentalManip	macro-F1	0.615	0.742
MultiManip	macro-F1	0.347	0.706
TalkDown	F1	0.058	0.729
ToxicChat	AUROC	0.587	0.587

was drawn from; we document the provenance in the data statement. Finally, the benchmark measures the exertion of agentive force, not its effect: a coercive dialogue that the listener successfully resists is still annotated as high-pressure, which is the correct label for detection but should not be read as evidence of harm.

J. Discussion

Three outcomes would matter beyond the specific numbers. First, if the dimension profiles are reliable, they provide a continuous, interpretable measurement of how language exerts pressure—covering transparent requests, exchange, normative pressure, and covert manipulation under one rubric—and locate each strategy family precisely in this space. Second, if each family exhibits a distinct dimension profile, system designers could select parsers and aggregation modes from a small set of validated configurations rather than re-tuning per task. Third, the seven-dimension profile offers an interpretable audit trail: a detection system can justify "this dialogue is coercive' with concrete dimension evidence (e.g., high D_3, low D_7), which is the kind of accountability that binary toxicity scores cannot provide. These implications motivate releasing not only the labels but also the gold dimension profiles and the parser outputs as part of the benchmark package.

VI. Conclusion

We presented LINGUAFORCE, a new dataset and unified benchmark for the agentive force of discourse in real-world multi-turn dialogues. The full release contains 3,432 dialogues annotated with dialogue-level intensity and seven universal psychological dimensions, and its reliability is verified through consistency with the existing coercion labels. We further specify three benchmark tasks, a two-stage dimension-aware pipeline, and research questions on dimension structure and zero-shot generalization. A zero-shot dimension-conditioned readout reaches binary AUC 0.852 and best F1 0.817 on the full release, exceeding a direct zero-shot judge, and transfers moderately to psychological-manipulation and patronizing-language corpora while remaining distinct from generic toxicity. Future work includes human inter-annotator agreement, cross-lingual extension, and studying how agentive force interacts with listener vulnerability and cultural norms.

References

[1] Z. Lin, Z. Wang, Y. Tong, Y. Wang, Y. Guo, Y. Wang, and J. Shang, "Toxicchat: Unveiling hidden challenges of toxicity detection in real-world user-ai conversation," in Findings of the Association for Computational Linguistics: EMNLP 2023, Singapore, 2023, pp. 4694–4702.

[2] Z. Wang and C. Potts, "Talkdown: A corpus for condescension detection in context," in Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), Hong Kong, China, 2019, pp. 3711–3719.

[3] J. Piskorski, N. Stefanovitch, G. D. S. Martino, and P. Nakov, "Semeval-2023 task 3: Detecting the category, the framing, and the persuasion techniques in online news in a multi-lingual setup," in Proceedings of the 17th International Workshop on Semantic Evaluation (SemEval-2023), Toronto, Canada, July 2023, pp. 2343–2361.

[4] Y. Wang, I. Yang, S. Hassanpour, and S. Vosoughi, "Mentalmanip: A dataset for fine-grained analysis of mental manipulation in conversations," in Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), Bangkok, Thailand, 2024, pp. 3747–3764.

[5] D. Khanna, P. Seth, S. Murali, A. Guru, S. Shukla, T. Tyagi, S. Chaurasia, and K. Ghosh, "Self-percept: Introspection improves large language models' detection of multi-person mental manipulation in conversations," in Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers), Vienna, Austria, 2025, pp. 660–675.

[6] Y. Gao, P. Gao, H. Bao, B. Li, J. Luo, Z. Wang, and W. Chen, "Boosting large language models for mental manipulation detection via data augmentation and distillation," arXiv preprint arXiv:2505.15255, 2025.

[7] G. Marwell and D. R. Schmitt, "Dimensions of compliance-gaining behavior: An empirical analysis," Sociometry, vol. 30, no. 4, pp. 350–364, 1967.

[8] R. L. Wiseman and W. J. Schenck-Hamlin, "A multidimensional scaling validation of an inductively-derived set of compliance-gaining strategies," Communication Monographs, vol. 48, no. 4, pp. 251–270, 1981.

[9] R. B. Cialdini, Influence: The Psychology of Persuasion. William Morrow, 1984.

[10] Anonymous, "Coercion: Benchmarking moral coercion in dialogues via psychological dimension decoupling," Under review, 2025, code and data: https://anonymous.4open.science/r/BMCDPDD-88D4/.

[11] J. L. Austin, How to Do Things with Words. Oxford University Press, 1962.

[12] J. R. Searle, Speech Acts: An Essay in the Philosophy of Language. Cambridge University Press, 1969.

[13] L. Breitfeller, E. Ahn, D. Jurgens, and Y. Tsvetkov, "Finding microaggressions in the wild: A case for locating elusive phenomena in social media posts," in Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), Hong Kong, China, 2019, pp. 1664–1674.

[14] J. Graham, J. Haidt, and B. A. Nosek, "Liberals and conservatives rely on different sets of moral foundations," Journal of Personality and Social Psychology, vol. 96, no. 5, pp. 1029–1046, 2009.

[15] K. Johnson and D. Goldwasser, "Classification of moral foundations in microblog political discourse," in Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), Melbourne, Australia, 2018, pp. 720–730.

[16] L. M. Ahearn, "Language and agency," Annual Review of Anthropology, vol. 30, pp. 109–137, 2001.

[17] T. Gebru, J. Morgenstern, B. Vecchione, J. W. Vaughan, H. Wallach, H. D. Iii, and K. Crawford, "Datasheets for datasets," vol. 64, no. 12, 2021, pp. 86–92.

[18] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. Chi, Q. Le, and D. Zhou, "Chain-of-thought prompting elicits reasoning in large language models," in Advances in Neural Information Processing Systems (NeurIPS), vol. 35, 2022.